

Example Run

Yagmur Onay

August 7, 2022

1 Train and deploy a model in Azure Machine Learning

1.1 Import Data

- Download the MNIST dataset
- Display some sample images

```
[1]: import os
from azureml.opendatasets import MNIST

data_folder = os.path.join(os.getcwd(), "/tmp/qs_data")
os.makedirs(data_folder, exist_ok=True)

mnist_file_dataset = MNIST.get_file_dataset()
mnist_file_dataset.download(data_folder, overwrite=True)

[1]: ['/tmp/qs_data/https%3A%2Fazureopendatastorage.azurefd.net/mnist/t10k-images-idx3-ubyte.gz',
      '/tmp/qs_data/https%3A%2Fazureopendatastorage.azurefd.net/mnist/t10k-labels-idx1-ubyte.gz',
      '/tmp/qs_data/https%3A%2Fazureopendatastorage.azurefd.net/mnist/train-images-idx3-ubyte.gz',
      '/tmp/qs_data/https%3A%2Fazureopendatastorage.azurefd.net/mnist/train-labels-idx1-ubyte.gz']
```

1.1.1 Take a look at the data

```
[2]: from utils import load_data
import matplotlib.pyplot as plt
import numpy as np
import glob

# note we also shrink the intensity values (X) from 0-255 to 0-1. This helps the
→model converge faster.
X_train = (
    load_data(
        glob.glob(
```

```

        os.path.join(data_folder, "**/train-images-idx3-ubyte.gz"),
↪recursive=True
    )[0],
    False,
)
    / 255.0
)
X_test = (
    load_data(
        glob.glob(
            os.path.join(data_folder, "**/t10k-images-idx3-ubyte.gz"),
↪recursive=True
        )[0],
        False,
    )
    / 255.0
)
y_train = load_data(
    glob.glob(
        os.path.join(data_folder, "**/train-labels-idx1-ubyte.gz"),
↪recursive=True
    )[0],
    True,
).reshape(-1)
y_test = load_data(
    glob.glob(
        os.path.join(data_folder, "**/t10k-labels-idx1-ubyte.gz"), recursive=True
    )[0],
    True,
).reshape(-1)

# now let's show some randomly chosen images from the training set.
count = 0
sample_size = 30
plt.figure(figsize=(16, 6))
for i in np.random.permutation(X_train.shape[0])[:sample_size]:
    count = count + 1
    plt.subplot(1, sample_size, count)
    plt.axhline("")
    plt.axvline("")
    plt.text(x=10, y=-10, s=y_train[i], fontsize=18)
    plt.imshow(X_train[i].reshape(28, 28), cmap=plt.cm.Greys)
plt.show()

```

1.2 Train model and log metrics with MLflow

Train the `LogisticRegression` classifier from the [SciKit Learn framework](#) to classify the data.

```
[3]: # create the model
import mlflow
import numpy as np
from sklearn.linear_model import LogisticRegression
from azureml.core import Workspace

# connect to your workspace
ws = Workspace.from_config()

# create experiment and start logging to a new run in the experiment
experiment_name = "azure-ml-in10-mins-tutorial"

# set up MLflow to track the metrics
mlflow.set_tracking_uri(ws.get_mlflow_tracking_uri())
mlflow.set_experiment(experiment_name)
mlflow.autolog()

# set up the Logistic regression model
reg = 0.5
clf = LogisticRegression(
    C=1.0 / reg, solver="liblinear", multi_class="auto", random_state=42
)

# train the model
with mlflow.start_run() as run:
    clf.fit(X_train, y_train)
```

```
2022/08/06 17:06:18 INFO mlflow.tracking.fluent: Autologging successfully
enabled for sklearn.
2022/08/06 17:06:18 INFO mlflow.tracking.fluent: Autologging successfully
enabled for pyspark.
2022/08/06 17:06:18 INFO mlflow.pyspark.ml: No SparkSession detected.
Autologging will log pyspark.ml models contained in the default allowlist. To
specify a custom allowlist, initialize a SparkSession prior to calling
mlflow.pyspark.ml.autolog() and specify the path to your allowlist file via the
spark.mlflow.pysparkml.autolog.logModelAllowlistFile conf.
2022/08/06 17:06:18 INFO mlflow.tracking.fluent: Autologging successfully
enabled for pyspark.ml.
```

1.3 View Experiment

In the left-hand menu in Azure Machine Learning Studio, select **Experiments** and then select your experiment.

1.4 Version control your models with the model registry

Once you have executed the code cell below you will be able to see the model in the registry by selecting **Models** in the left-hand menu in Azure Machine Learning Studio. Registered models are identified by name and version. Each time you register a model with the same name as an existing one, the registry increments the version.

```
[4]: # register the model
model_uri = "runs:{}/model".format(run.info.run_id)
model = mlflow.register_model(model_uri, "sklearn_mnist_model")
```

Registered model 'sklearn_mnist_model' already exists. Creating a new version of this model...

2022/08/06 17:09:00 INFO mlflow.tracking._model_registry.client: Waiting up to 300 seconds for model version to finish creation. Model

name: sklearn_mnist_model, version 4

Created version '4' of model 'sklearn_mnist_model'.

1.5 Deploy the model for real-time inference

Deploy a model so that an application can consume (inference) the model over REST.

1.5.1 Create deployment configuration

```
[5]: # create environment for the deploy
from azureml.core.environment import Environment
from azureml.core.conda_dependencies import CondaDependencies
from azureml.core.webservice import AciWebservice

# get a curated environment
env = Environment.get(
    workspace=ws,
    name="AzureML-sklearn-1.0-ubuntu20.04-py38-cpu",
    version=1
)
env.inferencing_stack_version='latest'

# create deployment config i.e. compute resources
aciconfig = AciWebservice.deploy_configuration(
    cpu_cores=1,
    memory_gb=1,
    tags={"data": "MNIST", "method": "sklearn"},
    description="Predict MNIST with sklearn",
)
```

1.5.2 Deploy model

Deploy the model to Azure Container Instance (ACI).

```
[6]: %%time
import uuid
from azureml.core.model import InferenceConfig
from azureml.core.environment import Environment
from azureml.core.model import Model

# get the registered model
model = Model(ws, "sklearn_mnist_model")

# create an inference config i.e. the scoring script and environment
inference_config = InferenceConfig(entry_script="score.py", environment=env)

# deploy the service
service_name = "sklearn-mnist-svc-" + str(uuid.uuid4())[:4]
service = Model.deploy(
    workspace=ws,
    name=service_name,
    models=[model],
    inference_config=inference_config,
    deployment_config=aciconfig,
)

service.wait_for_deployment(show_output=True)
```

Tips: You can try `get_logs()`: <https://aka.ms/debugimage#dockerlog> or local deployment: <https://aka.ms/debugimage#debug-locally> to debug if deployment takes longer than 10 minutes.

Running

```
2022-08-06 17:09:04+00:00 Creating Container Registry if not exists.
2022-08-06 17:09:05+00:00 Registering the environment.
2022-08-06 17:09:05+00:00 Use the existing image.
2022-08-06 17:09:05+00:00 Generating deployment configuration.
2022-08-06 17:09:06+00:00 Submitting deployment to compute.
2022-08-06 17:09:13+00:00 Checking the status of deployment sklearn-mnist-
svc-53b4..
2022-08-06 17:12:18+00:00 Checking the status of inference endpoint sklearn-
mnist-svc-53b4.
```

Succeeded

ACI service creation operation finished, operation "Succeeded"

CPU times: user 1.13 s, sys: 134 ms, total: 1.26 s

Wall time: 4min 38s

1.6 Test the model service

Test the model by sending a raw HTTP request to test the web service.

```
[7]: # send raw HTTP request to test the web service.
import requests
```

```

# send a random row from the test set to score
random_index = np.random.randint(0, len(X_test) - 1)
input_data = '{"data": [' + str(list(X_test[random_index])) + "]}'"

headers = {"Content-Type": "application/json"}

resp = requests.post(service.scoring_uri, input_data, headers=headers)

print("POST to url", service.scoring_uri)
print("label:", y_test[random_index])
print("prediction:", resp.text)

```

POST to url

http://7aa2b8fb-61ec-4f78-b9b7-71986888ff2b.westeurope.azurecontainer.io/score

label: 4

prediction: [4]

1.7 Clean up resources

If you want to keep workspace and only delete endpoint, delete the Model service using:

```
[8]: service.delete()
```